

NiftiArray: Solving Your Big Nifti Needs

source at https://github.com/muschellij2/niftiarray_talk

John Muschelli, slides at http://johnmuschelli.com/niftiarray_talk

Start of a Package: Plotting a voxel
from 800 Subjects in a Shiny App

How can you read in a NIfTI image
in R?

Options for reading in NIfTI images

Array like

- `oro.nifti::readNIfTI/neurobase::readnii`
- `AnalyzeFMRI::f.read.volume()` - can't read gzipped
- `neuroim::loadVolume("example.nii")`

C++ pointer

- `ANTsCore::antsImageRead`
- `RNifti::readNifti`

From RNifti:

<https://github.com/jonclayden/RNifti#performance>

```
library(microbenchmark)
microbenchmark(AnalyzeFMRI::f.read.volume("example.nii"),
               ANTsRCore::antsImageRead("example.nii"),
               neuroim::loadVolume("example.nii"),
               oro.nifti::readNIFTI("example.nii"),
               RNifti::readNifti("example.nii"),
               tractor.base::readImageFile("example.nii"), unit="ms")
# Unit: milliseconds
#           expr      min       lq      mean
# AnalyzeFMRI::f.read.volume("example.nii") 26.312881 26.769244 29.685981
#   ANTsRCore::antsImageRead("example.nii")  1.150787  1.626918  2.149145
#       neuroim::loadVolume("example.nii") 34.596506 37.732245 54.065227
#   oro.nifti::readNIFTI("example.nii") 57.059828 61.953430 89.386706
#       RNifti::readNifti("example.nii")  0.986773  1.136562  1.683821
# tractor.base::readImageFile("example.nii") 33.380407 34.096574 34.961812
#   median      uq      max neval
# 27.301693 28.214908 185.441664   100
#  2.210696  2.481675   3.376350   100
# 40.714061 45.425047 192.978225   100
# 65.064212 71.425561 220.709246   100
#  1.501528  1.856601   7.961566   100
# 34.617259 35.257633 42.108584   100
```

What do we want

1. Something fast
2. Has overloaded functions (e.g. `image + image = image`)
3. Can use most plotting functions on it (e.g. `ortho2`)
4. Is array like but has header
5. Keeps memory low
6. Has **random access**

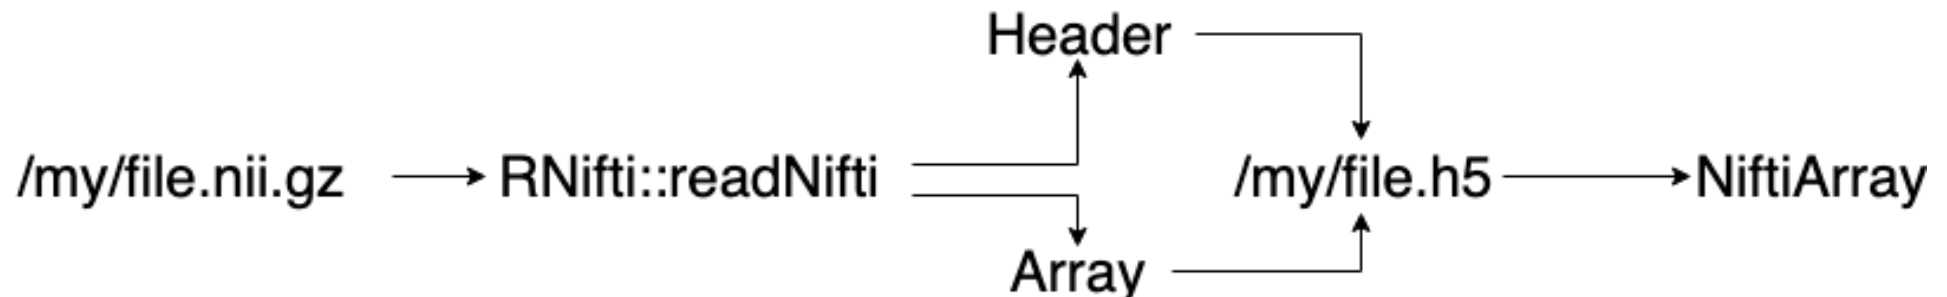
Current issues

- `oro.nifti::readNifti` - slow
- `ANTsCore::antsImageRead` - installation is
- `RNifti::readNifti` - from <https://github.com/jonclayden/RNifti/issues/14> :

For your single-precision floating-point image with about 1bn elements, the “native” data size in the C struct is roughly 4 GiB, while the R array uses 8 GiB because it uses double precision. When you use `readNifti(..., internal=TRUE)` the resulting object has only the C struct, but with `readNifti(..., internal=FALSE)` the object has both the C struct and the R array, for a total of 12 GiB.

Solution: NiftiArray

- Stores image and header in HDF5 (Hierarchical Data Format) format
- Based on `HDF5Array` from Bioconductor (<https://bioconductor.org/packages/release/bioc/html/HDF5Array.html>)



Takes longer than RNifti (the first time)

```
fname = fslr::mni_fname()
system.time({
  rnifti = RNifti::readNifti(fname)
})
```

```
##      user  system elapsed
##    0.140    0.021    0.296
```

```
system.time({
  narray = NiftiArray::writeNiftiArray(fname)
  narray@seed@filepath
  narray
})
```

```
##      user  system elapsed
##    4.721    0.133    5.450
```

Sizes

```
format(object.size(rnifti), units = "Kb")
```

```
## [1] "28208.3 Kb"
```

```
format(object.size(narray), units = "Kb")
```

```
## [1] "8.7 Kb"
```

Header Information

```
print(nifti_header(narray))
```

```
## NIfTI-1 header
##      sizeof_hdr: 348
##      dim_info: 0
##      dim: 3  182  218  182  1  1  1  1
##      intent_p1: 0
##      intent_p2: 0
##      intent_p3: 0
##      intent_code: 0
##      datatype: 8
##      bitpix: 32
##      slice_start: 0
##      pixdim: -1  1  1  1  0  0  0  0
##      vox_offset: 352
##      scl_slope: 0
##      scl_inter: 0
##      slice_end: 0
##      slice_code: 0
##      xyzt_units: 10
##      cal_max: 8000
##      cal_min: 3000
```

Accessing voxels

```
narray[18, 5, 14]
```

```
## [1] 572
```

```
sum(narray > 9990)
```

```
## [1] 20
```

```
narray[ narray > 9990]
```

```
## [1] 9999 9999 9992 9997 9991 9993 9997 9998 9999 9999 9994 9997 9998 9999
```

```
## [15] 9998 9991 9993 9991 9996 9994
```

```
DelayedArray::extract_array(narray, index = list(NULL, NULL, 1))
```

```
## , , 1
```

```
##
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12]
```

```
## [1,]    0    0    0    0   30   31   29   29   29   29   29   28
```

What are the benefits?

- Overall most benefits come from saving H5 files on “big” computer
- Random access - the `DelayedArray` package
- Operations on `DelayedArray` objects, the `DelayedMatrixStats` package
<https://bioconductor.org/packages/release/bioc/html/DelayedMatrixStats.html>

What are the benefits? Reshaping to a Matrix

```
d = dim(narray)
d
```

```
## [1] 182 218 182
```

```
dim(narray) = c(d, 1)
new_dim = c(prod(d), 1)
newarray = writeNiftiArray(narray)
mat = ReshapedNiftiArray(newarray@seed@filepath, dim = new_dim)
dim(mat)
```

```
## [1] 7221032      1
```

```
head(mat)
```

```
## <6 x 1> DelayedMatrix object of type "integer":
##      [,1]
## [1,]    0
## [2,]    0
## [3,]    0
```

Why? Voxel-wise operations in low-memory setting

```
big_mat = DelayedArray::acbind(mat, mat, mat)
dim(big_mat)
```

```
## [1] 7221032      3
```

```
format(object.size(big_mat), units = "Kb")
```

```
## [1] "26.7 Kb"
```

```
rmed = DelayedMatrixStats::rowMedians(big_mat)
```

Why? 4D Example

```
nii_fname = system.file("extdata",  
                          "example_4d.nii.gz", package = "RNifti")  
from = NiftiArray::NiftiArray(nii_fname)  
dim(from)
```

```
## [1] 96 96 60 3
```

```
A <- ReshapedNiftiArray(filepath = from@seed@filepath,  
                        name = from@seed@name,  
                        dim = c(prod(dim(from)[1:3]), dim(from)[4]))  
dim(A)
```

```
## [1] 552960 3
```

```
med = DelayedMatrixStats::rowMedians(A)
```


Back to an image

```
arr = array(med, dim = dim(from)[1:3])  
med_array = writeNiftiArray(arr, header = nifti_header(from))  
as(med_array, "niftiImage")
```

```
## Image array of mode "double" (4.2 Mb)  
## - 96 x 96 x 60 voxels  
## - 2.5 x 2.5 x 2.5 mm per voxel
```

Issues

- Need to write H5 file first (do in parallel on cluster)
- Memory tradeoff with speed due to chunking
- Most of this needs things in the same space

Use cases

- Something on Shiny for speed or memory
 - shinyapps.io has very restrictive memory settings for imaging
- Analyzing 1000s of subjects together
- Pass object from R to Matlab:
<https://www.mathworks.com/help/matlab/ref/hdf5read.html>

Extensions

- Anything that works with DelayedArrays should work here
- Check out Bioconductor
- Example: `mbkmeans` - mini-batch k-means
<https://bioconductor.org/packages/release/bioc/html/mbkmeans.html>